

Inside the RNN: update equations

An RNN layer, at time step t

- Takes input element $\mathbf{x}_{(t)}$
- Updates latent state $\mathbf{h}_{(t)}$
- Optionally outputs $\mathbf{y}_{(t)}$

according to the equations

$$\begin{aligned}\mathbf{h}_{(t)} &= \phi(\mathbf{W}_{xh}\mathbf{x}_{(t)} + \mathbf{W}_{hh}\mathbf{h}_{(t-1)} + \mathbf{b}_h) \\ \mathbf{y}_{(t)} &= \mathbf{W}_{hy}\mathbf{h}_{(t)} + \mathbf{b}_y\end{aligned}$$

where

- ϕ is an activation function (usually \tanh)
- \mathbf{W} are the weights of the RNN layer
 - partitioned into \mathbf{W}_{xh} , \mathbf{W}_{hh} , \mathbf{W}_{hy}
 - \mathbf{W}_{xh} : weights that update $\mathbf{h}_{(t)}$ based on $\mathbf{x}_{(t)}$
 - \mathbf{W}_{hh} : weights that update $\mathbf{h}_{(t)}$ based on $\mathbf{h}_{(t-1)}$
 - \mathbf{W}_{hy} : weights that update $\mathbf{y}_{(t)}$ based on $\mathbf{h}_{(t)}$

RNN

Notes

- The RNN literature uses ϕ rather than $a_{(l)}$ to denote an activation function
- This is the update equation for a single example $\mathbf{x}^{(i)}$
- In practice, we can simultaneously update for *multiple examples*
 - The $m' < m$ examples in a minibatch, as examples are independent

Let's try to understand these equations

$$\mathbf{h}_{(t)} = \phi(\mathbf{W}_{xh}\mathbf{x}_{(t)} + \mathbf{W}_{hh}\mathbf{h}_{(t-1)} + \mathbf{b}_h)$$

$\mathbf{h}_{(t)}$ is the latent state after time step t

- It is a *vector* of length $||\mathbf{h}||$
- We drop the time subscript as the dimension on each step is the same

$\mathbf{W}_{xh} \mathbf{x}_{(t)}$ must therefore *also be a vector* of length $||\mathbf{h}||$

- $||\mathbf{W}_{xh}||$ is a matrix of shape $(||\mathbf{h}|| \times ||\mathbf{x}||)$
- \mathbf{h}_j , the j^{th} element of latent state \mathbf{h} is the dot product of row j of \mathbf{W}_{xh} and \mathbf{x}
- So $\mathbf{W}_{xh}^{(j)}$ describes how input $\mathbf{x}_{(t)}$ influences new state $h_{(t),j}$

That is: there are separate weights for each j that describe the interaction of \mathbf{h} and \mathbf{x}

Similarly, $\mathbf{W}_{hh} \mathbf{h}_{(t-1)}$ must be a vector of length $\|\mathbf{h}\|$

- $\|\mathbf{W}_{hh}\|$ is a matrix of shape $(\|\mathbf{h}\| \times \|\mathbf{h}\|)$
- So $\mathbf{W}_{hh}^{(j)}$ describes how prior state $\mathbf{h}_{(t-1)}$ influences new state $\mathbf{h}_{(t),j}$

\mathbf{b}_h , the bias/threshold must also be a vector of length $||\mathbf{h}||$

- It adjusts the threshold of activation function ϕ
- As per our practice: we will usually fold \mathbf{b} into the weight matrices $\mathbf{W}_{xh}, \mathbf{W}_{hh}$

Finally, activation ϕ maps a vector of length $\|\mathbf{h}\|$ to another vector of length $\|\mathbf{h}\|$

- The updated state

So updated latent state $\mathbf{h}_{(t)}$ is influenced

- By the input $\mathbf{x}_{(t)}$
- The prior latent state $\mathbf{h}_{(t-1)}$

The second equation

$$\mathbf{y}_{(t)} = \mathbf{W}_{hy}\mathbf{h}_{(t)} + \mathbf{b}_y$$

is just a "translation" of the latent state $\mathbf{h}_{(t)}$

- To $\mathbf{y}_{(t)}$, the t^{th} element of the output sequence
- $||\mathbf{W}_{hy}||$ is a matrix of shape $(||\mathbf{y}|| \times ||\mathbf{h}||)$
 - $||\mathbf{y}||$ is the length of each output element and is problem dependent
 - For example: a OHE

It is common to equate $\mathbf{y}_{(t)} = \mathbf{h}_{(t)}$

- No separate "output"
- Just the latent state
- Particularly when using stacked RNN layers
 - $\mathbf{y}_{(t)}$ becomes the input to the next layer

Equation in pseudo-matrix form

You will often see a short-hand form of the equation.

Look at $\mathbf{h}_{(t)}$ as a function of two inputs $\mathbf{x}, \mathbf{h}_{(t-1)}$.

We can stack the two inputs into a single matrix.

Stack the two matrices $\mathbf{W}_{xh}, \mathbf{W}_{hh}$ into a single weight matrix

$$\mathbf{h}_{(t)} = \mathbf{W}\mathbf{I} + \mathbf{b}$$

with

$$\mathbf{W} = [\mathbf{W}_{xh} \quad \mathbf{W}_{hh}]$$

$$\mathbf{I} = \begin{bmatrix} \mathbf{x}_{(t)} \\ \mathbf{h}_{(t-1)} \end{bmatrix}$$

Stacked RNN layers revisited

With the benefit of the RNN update equations, we can clarify how stack RNN layers works.

Let superscript $[l]$ denote a stacked layer of RNN.

So the RNN update equation for the bottom layer 1 becomes

$$\mathbf{h}_{(t)}^{[1]} = \phi(\mathbf{W}_{xh}\mathbf{x}_{(t)} + \mathbf{W}_{hh}\mathbf{h}_{(t-1)}^{[1]} + \mathbf{b}_h)$$

The RNN update equation for layer $[l]$ becomes

$$\mathbf{h}_{(t)}^{[l]} = \phi(\mathbf{W}_{xh}\mathbf{h}_{(t)}^{[l-1]} + \mathbf{W}_{hh}\mathbf{h}_{(t-1)}^{[l]} + \mathbf{b}_h)$$

That is: the input to layer $[l]$ is $\mathbf{h}_{(t)}^{[l-1]}$ rather than $\mathbf{x}_{(t)}$

Loss function

As usual, the objective of training is to find the weights \mathbf{W} that minimize a loss function

$$\mathcal{L} = L(\hat{\mathbf{y}}, \mathbf{y}; \mathbf{W})$$

which is the average of per example losses $\mathcal{L}^{(i)}$

$$\mathcal{L} = \frac{1}{m} \sum_{i=1}^m \mathcal{L}^{(i)}$$

When the output is a sequence

- It's important to recognize that the *target* is a sequence too !
- So the per example loss has an added temporal dimension
- Loss per example *per time step*
- Comparing the *predicted* t^{th} output $\hat{\mathbf{y}}_{(t)}^{(i)}$ to the t^{th} target $\mathbf{y}_{(t)}^{(i)}$

In the case that the API outputs sequences

- $\mathcal{L}^{(\mathbf{i})} = \sum_{t=1}^T \mathcal{L}_{(t)}^{(\mathbf{i})}$

In the case that the API outputs a single value

- $\mathcal{L}^{(\mathbf{i})} = \mathcal{L}_{(T)}$

RNN Loss: Forward pass

In [2]: `print("Done")`

Done